

PER-CLASS THRESHOLD CALIBRATION FOR TRAINING-FREE OPEN-VOCABULARY SEMANTIC SEGMENTATION OF REMOTE SENSING IMAGERY

MID-PROJECT REVIEW

PRIYANSHU KUMAR RAJ GAUTAM
under the supervision of
Dr. Shounak Chakraborty & Prof. K Sathya Babu

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING,
INDIAN INSTITUTE OF INFORMATION TECHNOLOGY DESIGN AND MANUFACTURING, KURNOOL

September 10, 2026

PRESENTATION OUTLINE

- 1 Introduction
- 2 Motivation & Applications
- 3 Problem Statement
- 4 Literature Survey
- 5 Research Gaps
- 6 Objectives
- 7 Experimental Setup & Results
- 8 Application Study

1. INTRODUCTION

Open-vocabulary semantic segmentation of aerial imagery. The input is an image together with a list of class names given as plain words; the output is a class label for every pixel. The setting is **training-free**: no network weights are learned, and the class list may change at inference without retraining.

Baseline. SegEarth-OV3 [1] applies SAM 3 [2] to remote sensing. For each class prompt, SAM 3 emits a presence score, a dense semantic map and a set of instance masks; SegEarth-OV3 fuses these and applies a single confidence threshold τ , tuned per dataset. Pixels below τ are assigned to a catch-all class. **Reproduced on our hardware at 47.38 mIoU against a published 47.4.**

Contribution

A single confidence threshold is the wrong *shape* for a multi-class pipeline. We fit **one threshold per class**, and **one scale per class applied before the argmax**, using the same supervision budget the baseline already spends tuning its single τ . **No network weights are trained.**

2. MOTIVATION — THE BASELINE DISCARDS LAND COVER

Measured on LoveDA validation at SegEarth-OV3's own published $\tau = 0.5$:


- **29.68%** of all real-class pixels are assigned to the catch-all class **background** — **323 million pixels**.
- Discard exceeds real-class confusion by **3:1**. The dominant error mode is abstention rather than misclassification.
- **water**: precision **89.5%**, recall **54.7%**. The class is correct nine times in ten when predicted, yet recovers barely half of the water present.

Cost of global remedies

Lower τ to 0.1	−5.54 mIoU; 1.73 incorrect pixels per correct one
Remove presence gating	−11.97 mIoU
No intervention	29.68% of land cover discarded

A correction that operates uniformly across classes therefore cannot recover the residual. **The correction must be per-class.**

2. MOTIVATION — THE DISCARD, VISUALISED



A LoveDA validation tile. Left to right: **input image**, **ground truth**, **SegEarth-OV3 prediction at $\tau = 0.5$** , and the **pixels assigned to background**. Entire water regions are discarded — not misclassified, but declared unknown — because their confidence falls below a threshold tuned for the dataset as a whole rather than for **water**.

2. APPLICATION DOMAINS

The setting is intended for deployments in which the class list is not fixed in advance and labelled data for the target region does not exist.

- **Disaster response.** Flood or damage extent from UAV imagery acquired within hours of an event, where no annotated data exists for the affected district and there is no time to train a model.
- **Infrastructure monitoring.** Classes such as solar panel or transmission line appear in no land-cover benchmark, yet can be supplied to the pipeline as words.
- **Settlement mapping.** Vocabularies such as tin roof, unpaved road and water tank change between surveys and between regions.

Deployment cost

The proposed correction requires **~200 labelled tiles** and **no gradient steps**. Adapting the same pipeline by fine-tuning would require labelled data in the thousands of tiles, a training run, and would forfeit the training-free property.

3. PROBLEM STATEMENT — FORMULATION

Let $\mathcal{C} = \{1, \dots, N\}$ denote the vocabulary and $s_c(p) \in [0, 1]$ the score the pipeline assigns to class c at pixel p . Let b denote the catch-all class.

Baseline decision rule, with a single global threshold τ :

$$\hat{y}(p) = \begin{cases} \arg \max_c s_c(p), & \text{if } \max_c s_c(p) \geq \tau \\ b, & \text{otherwise} \end{cases}$$

Proposed rule, with parameters fitted on a calibration set of ~ 200 labelled tiles:

$$\hat{y}(p) = \begin{cases} c^*, & \text{if } s_{c^*}(p) \geq \tau_{c^*} \\ b, & \text{otherwise} \end{cases} \quad c^* = \arg \max_c \mathbf{w}_c s_c(p)$$

Objective

$(\mathbf{w}, \tau) = \arg \max \text{mIoU}_{\text{real}}$, optimised by coordinate ascent. **No network weights are modified.** The threshold is applied to the *raw* score, so $\mathbf{w}$ influences only the **argmax**.

3. COMPLETENESS OF THE PER-CLASS THRESHOLD FAMILY

Proposition. For a fixed `argmax`, any strictly increasing per-class recalibration g_c composed with a global threshold τ is equivalent to a per-class threshold vector:

$$g_c(s) \geq \tau \iff s \geq g_c^{-1}(\tau) = \tau_c.$$

- **Per-class thresholds therefore exhaust every correction reachable by rescaling scores after the `argmax`**; no member of that family can outperform them.
- Scaling applied *before* the `argmax` alters which class wins, and so constitutes a **strictly larger family**.

Empirical consequence

On LoveDA, **6.0%** of catch-all assignments are pixels for which the catch-all class won the `argmax` at $s \geq \tau$. A threshold cannot alter an `argmax`, so this residual is unreachable by the first parameter set and motivates the second.

4. LITERATURE SURVEY

No.	Work	Year	Contribution and relevance
1	CLIP [3]	2021	Image–text contrastive pretraining; origin of the open-vocabulary paradigm.
2	MaskCLIP [4]	2022	First extraction of dense predictions from CLIP by removing pooling.
3	GroupViT [5]	2022	Learns segment grouping from text supervision alone.
4	SCLIP [6]	2023	Correlative self-attention; the codebase from which SegEarth-OV3 derives.
5	SAM [7]	2023	Promptable class-agnostic segmentation without a fixed label set.
6	ClearCLIP [8]	2024	Identifies which CLIP components degrade dense prediction.
7	ProxyCLIP [9]	2024	Self-supervised features as structural guidance for segmentation.
8	SegEarth-OV [10]	2025	Feature upsampling and global bias subtraction for remote sensing; establishes the evaluation protocol.
9	SAM 3 [2]	2025	Presence, semantic and instance heads unified in one promptable model.
10	SegEarth-OV3 [1]	2025	SAM 3 for remote sensing: dual-head fusion and presence gating. Baseline of this work.
11	ConInfer [11]	2026	Gaussian mixture over self-supervised patch features fused

5. RESEARCH GAPS

Gap 1 — the decision rule is unexamined

Every method in this family terminates in an `argmax` followed by a single global threshold, tuned per dataset. Research effort is directed at improving the *scores*; whether the *decision rule* applied to them has the appropriate form has not been studied.

Gap 2 — the discarded residual is not quantified

Published work reports mIoU. It does not report what fraction of real land cover is assigned to the catch-all class, nor whether that fraction is recoverable.

Gap 3 — catch-all classes distort the reported metric

mIoU is an *unweighted* mean over classes, so a catch-all class contributes $1/N$ of the score irrespective of its semantic content. A method may therefore appear to improve or degrade segmentation while land-cover quality is unchanged.

6. OBJECTIVES

Completed

- 1 **Reproduce** SegEarth-OV3 on local hardware, so that every subsequent measurement is made against a baseline we control. 47.38 vs published 47.4
- 2 **Quantify** the residual the baseline discards, and determine whether any global adjustment recovers it. 29.68%; no global adjustment does
- 3 **Design** a correction that trains no weights and consumes only the supervision budget the baseline already spends. two parameter sets

In progress

- 4 **Extend** the correction beyond a single vector per dataset: the fitted parameters are currently constant across tiles, and a scene-adaptive formulation remains to be investigated.
- 5 **Investigate** a per-class fusion of SAM 3's semantic and instance heads, which are presently combined by a fixed maximum.
- 6 **Broaden** the evaluation to further high-resolution benchmarks and establish the conditions under which the correction transfers.

7. EXPERIMENTAL SETUP

Datasets

	Tiles	GSD	Classes
LoveDA (val)	1669	0.3 m	7
ISPRS Potsdam	2016	5 cm	6
OpenEarthMap	384	0.25–0.5 m	9

Baselines and controls

- SegEarth-OV3, reproduced — the primary baseline
- Global τ sweep; presence-gating removal
- Per-class Otsu and eleven further label-free rules

Protocol

- Five-fold cross-validation. Calibration and evaluation tiles are **always disjoint**.
- The baseline is recomputed on the *same* held-out tiles, so both rules are scored on identical pixels.
- Both metrics are reported: full mIoU and mIoU excluding the catch-all class.
- Results are confirmed by the official evaluator, not by arithmetic on a cached statistic.

Validation gate


Every instrumented run must still reproduce **47.37 mIoU** and **29.68%** discard. A deviation indicates that the instrumentation altered model behaviour.

7. RESULTS

	LoveDA 1669 tiles	Potsdam 2016 tiles	OpenEarthMap 384 tiles
SegEarth-OV3, reproduced	47.38	57.60	44.16
+ per-class thresholds	$+1.18 \pm 0.45$	+0.59	+0.30
+ per-class scaling	$+1.16 \pm 0.19$	$+4.86 \pm 0.35$	—
Total Δ mIoU	+2.32	+5.67	+0.30
Δ excluding catch-all	+1.36	+6.00	+1.75
Folds positive	5 / 5	5 / 5	—

- **Confirmed end to end.** On Potsdam the official evaluator gives $57.60 \rightarrow 58.35 \rightarrow 63.27$ over 1816 held-out tiles, each stage within **0.05** mIoU of its predicted value.
- **Largest per-class change:** Potsdam tree, **37.92 $\rightarrow$ 59.09** IoU — a class with precision 93.34% and recall 38.63%, the asymmetry the method is designed to correct.
- **OpenEarthMap** is positive under the catch-all-excluded metric (+1.75) although full mIoU appears flat: its catch-all class contributes one ninth of the mean and offsets the gain.

7. RESULTS — QUALITATIVE



Four ISPRS Potsdam tiles. Columns: input, ground truth, SegEarth-OV3, and the calibrated rule. **The network, its weights and the forward pass are identical across the two right-hand columns; only the decision rule differs.** The fourth row is a tile on which the calibrated rule **loses 7.72 mIoU**: the


8. APPLICATION STUDY — AN UNSEEN REGION AND VOCABULARY

The correction was applied to 25 tiles cut from 5 openly licensed UAV scenes over India at 4–5 cm ground sample distance — a region present in none of the three benchmarks — under two class lists supplied as words at inference.

- **The segmentation is stable under renaming.** A flooded building is partitioned almost identically by both vocabularies: 78.6% **water** against 79.1% **flood water**, and 18.2% **concrete roof** against 18.3% **building**. The regions are recovered consistently; only the labels follow the supplied words.
- **A class absent from the vocabulary is not recovered.** 100.0% of a photovoltaic-array tile is assigned to a single class until **solar panel** is supplied, which then accounts for 4.6% — below the visible extent of the panels.

No ground truth exists for this imagery. The quantities above are per-class pixel shares and are not accuracy measurements.

8. APPLICATION STUDY — ONE SCENE, TWO VOCABULARIES



Indian UAV imagery at 4–5 cm. Identical weights and forward pass; only the supplied class list differs between the second and third columns. **Lower row:** a building in flood water, recovered consistently under both vocabularies. **Upper row:** a photovoltaic array, for which the first vocabulary provides no

PROJECT TIMELINE

Period	Activity	Status
Aug 2026	Environment, baseline reproduction, discard measurement	Completed
Aug 2026	Co-occurrence prior: constructed, evaluated, rejected	Completed
Aug–Sep 2026	Per-class thresholds; cross-validation and deployment	Completed
Sep 2026	Per-class scaling; end-to-end verification on two datasets	Completed
Sep 2026	Application study on unseen imagery	Completed
Oct–Dec 2026	Scene-adaptive formulation; head-fusion investigation	Ongoing
Oct–Dec 2026	Additional high-resolution benchmarks	Ongoing
Jan–Feb 2027	Manuscript preparation and internal review	Planned
Q1 2027	Journal submission	Planned

REFERENCES I



K. Li, S. Zhang, Y. Wang, Y. Deng, Z. Wang, D. Meng, and X. Cao, “SegEarth-OV3: Exploring SAM 3 for open-vocabulary semantic segmentation in remote sensing images,” *arXiv preprint arXiv:2512.08730*, 2025. v2, 22 Apr. 2026.



N. Carion, L. Gustafson, Y.-T. Hu, S. Debnath, R. Hu, D. Suris, *et al.*, “SAM 3: Segment anything with concepts,” *arXiv preprint arXiv:2511.16719*, 2025.



A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, “Learning transferable visual models from natural language supervision,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2021.



C. Zhou, C. C. Loy, and B. Dai, “Extract free dense labels from CLIP,” in *Proc. European Conf. Computer Vision (ECCV)*, 2022.



J. Xu, S. De Mello, S. Liu, W. Byeon, T. Breuel, J. Kautz, and X. Wang, “GroupViT: Semantic segmentation emerges from text supervision,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2022.



F. Wang, J. Mei, and A. Yuille, “SCLIP: Rethinking self-attention for dense vision–language inference,” in *Proc. European Conf. Computer Vision (ECCV)*, 2024.

REFERENCES II



A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, P. Dollár, and R. Girshick, “Segment anything,” in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2023.



M. Lan, C. Chen, Y. Ke, X. Wang, L. Feng, and W. Zhang, “ClearCLIP: Decomposing CLIP representations for dense vision–language inference,” in *Proc. European Conf. Computer Vision (ECCV)*, 2024.



M. Lan, C. Chen, Y. Ke, X. Wang, L. Feng, and W. Zhang, “ProxyCLIP: Proxy attention improves CLIP for open-vocabulary segmentation,” in *Proc. European Conf. Computer Vision (ECCV)*, 2024.



K. Li, R. Liu, X. Cao, X. Bai, F. Zhou, D. Meng, and Z. Wang, “SegEarth-OV: Towards training-free open-vocabulary segmentation for remote sensing images,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2025.



W. Chen, Z. Hu, Y. Zhang, H. Ning, and Y. Tai, “ConInfer: Context-aware inference for training-free open-vocabulary remote sensing segmentation,” *arXiv preprint arXiv:2603.29271*, 2026.



C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2017.



N. Otsu, “A threshold selection method from gray-level histograms,” *IEEE Trans. Systems, Man, and Cybernetics*, vol. 9, no. 1, pp. 62–66, 1979.

REFERENCES III



P. Krähenbühl and V. Koltun, “Efficient inference in fully connected CRFs with gaussian edge potentials,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2011.

Thank You